TRANSFORMER MODELS FOR MOTION PATTERN ANALYSIS

ARTIFICIAL INTELLIGENCE & MACHINE LEARNING

PROJECT REPORT

INDIAN INSTITUTE OF TECHNOLOGY PATNA

DEPARTMENT OF ELECTRICAL ENGINEERING

SUBMITTED BY

ANKIT KUMAR KERO (2302CM07)

ARAVIND (2301EC41)

SANTOSH (2301EC24)

PRITHVI (2301EC36)

Academic Year: 2025–26

Abstract

Human motion prediction is a critical problem in modern artificial intelligence and computer vision.
It involves forecasting articulated human movements from historical pose sequences. Traditional approaches
such as Recurrent Neural Networks and Long Short-Term Memory models suffer from limitations in capturing
long-range temporal dependencies. Transformer architectures address these limitations by leveraging
self-attention mechanisms that model global spatial and temporal relationships. This work aims to build a
scalable and robust framework for long-horizon human motion prediction with improved stability and realism.
The applications of this research include robotics, healthcare, intelligent monitoring, and immersive digital media

🔷 **Keywords**

Keywords:
Transformer, Motion Prediction, Self-Attention, Human Pose Estimation, Deep Learning, Sequence Modeling, MPJPE

1. Introduction

Human motion pattern analysis is an essential area of research in artificial intelligence, computer vision, and biomechanics. It focuses on understanding and predicting human movements based on observed motion sequences. Accurate motion prediction plays a crucial role in various real-world applications such as robotics, healthcare monitoring, autonomous systems, sports analytics, and virtual reality environments.

The complexity of human motion arises from nonlinear dynamics, temporal dependencies, and variations in human behavior. Traditional sequence modeling techniques, including Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) models, have been widely used for motion prediction. However, these models process data sequentially, making them

computationally expensive and less effective in capturing long-range dependencies.

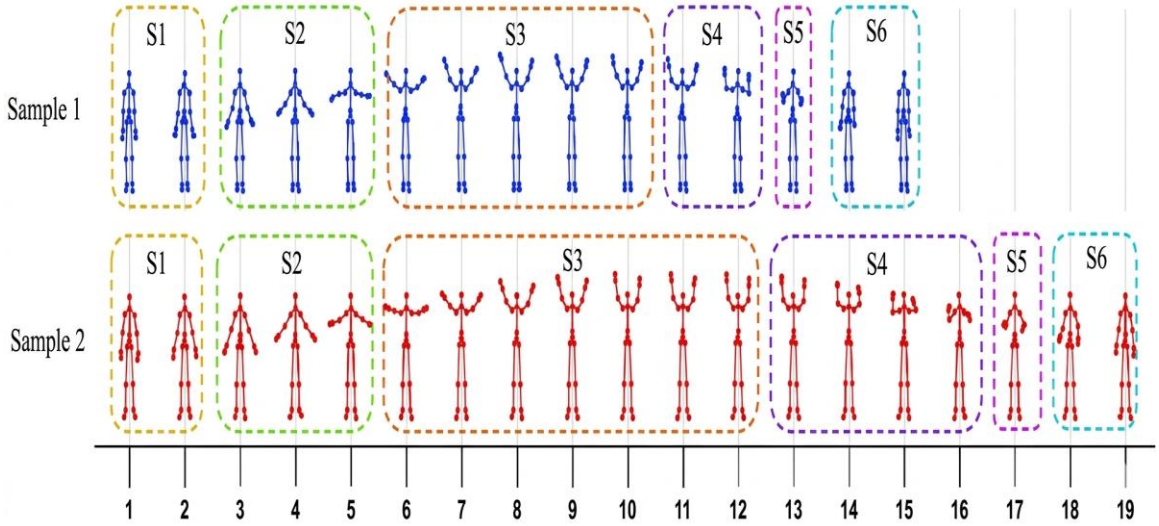


Figure 1: proposed Motion Pattern Analysis system, illustrating the transition from raw skeletal

Recently, Transformer-based architectures have emerged as a powerful alternative due to their ability to process sequences in parallel and capture global dependencies using self-attention mechanisms. Unlike RNNs, Transformers can model both spatial and temporal relationships efficiently, making them highly suitable for human motion prediction tasks.

In this work, we propose a Transformer-based framework for motion pattern analysis, which takes historical pose sequences as input and predicts future motion trajectories. The model aims to improve prediction accuracy, stability, and scalability for long-term motion forecasting.

1.1 Applications

Applications include autonomous robotics, healthcare monitoring, rehabilitation, augmented and virtual reality,
sports analytics, and smart surveillance. These systems require accurate

prediction of human trajectories
to enhance safety, efficiency, and realism.

2. Motivation

Earlier sequence modeling relied on statistical models and recurrent neural networks. Although LSTMs
improved performance, they remain sequential and computationally expensive. The vanishing gradient problem
limits their ability to learn long-term dependencies. Transformers overcome these issues through parallel
processing and attention-based global context modeling.

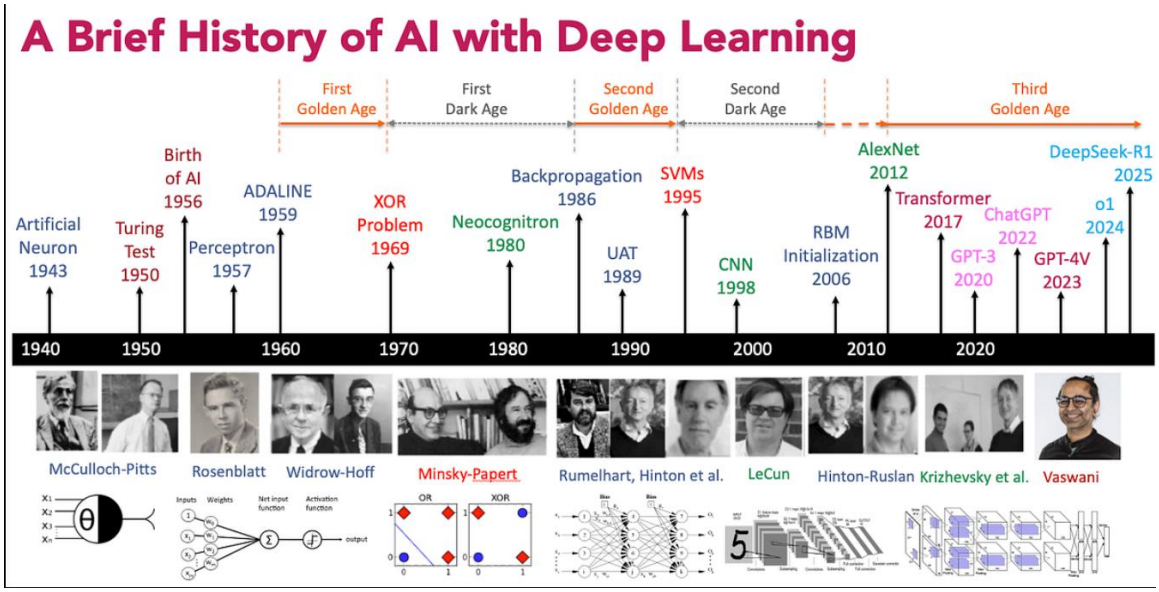
3. Problem Statement

The objective is to design a Transformer-based sequence-to-sequence framework. Given observed pose
sequences, the system predicts future human motion while preserving spatial consistency and temporal smoothness.
The model must generate physically realistic motion trajectories that align with biomechanical constraints.

4. Literature Review

Human motion prediction has been extensively studied using various sequence modeling techniques. Early approaches relied on statistical models and basic machine learning methods, which were limited in capturing complex temporal dynamics.

Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) networks became popular for motion prediction tasks due to their ability to model sequential data. However, these models suffer from issues such as vanishing gradients, limited parallelization, and difficulty in capturing long-range dependencies.



Recent advancements introduced attention-based models and Transformer architectures, which have shown significant improvements in sequence modeling tasks. Transformers eliminate recurrence and instead rely entirely on self-attention mechanisms to capture global dependencies efficiently. According to the "Attention Is All You Need" paper, Transformers enable parallel computation and achieve superior performance compared to traditional models .

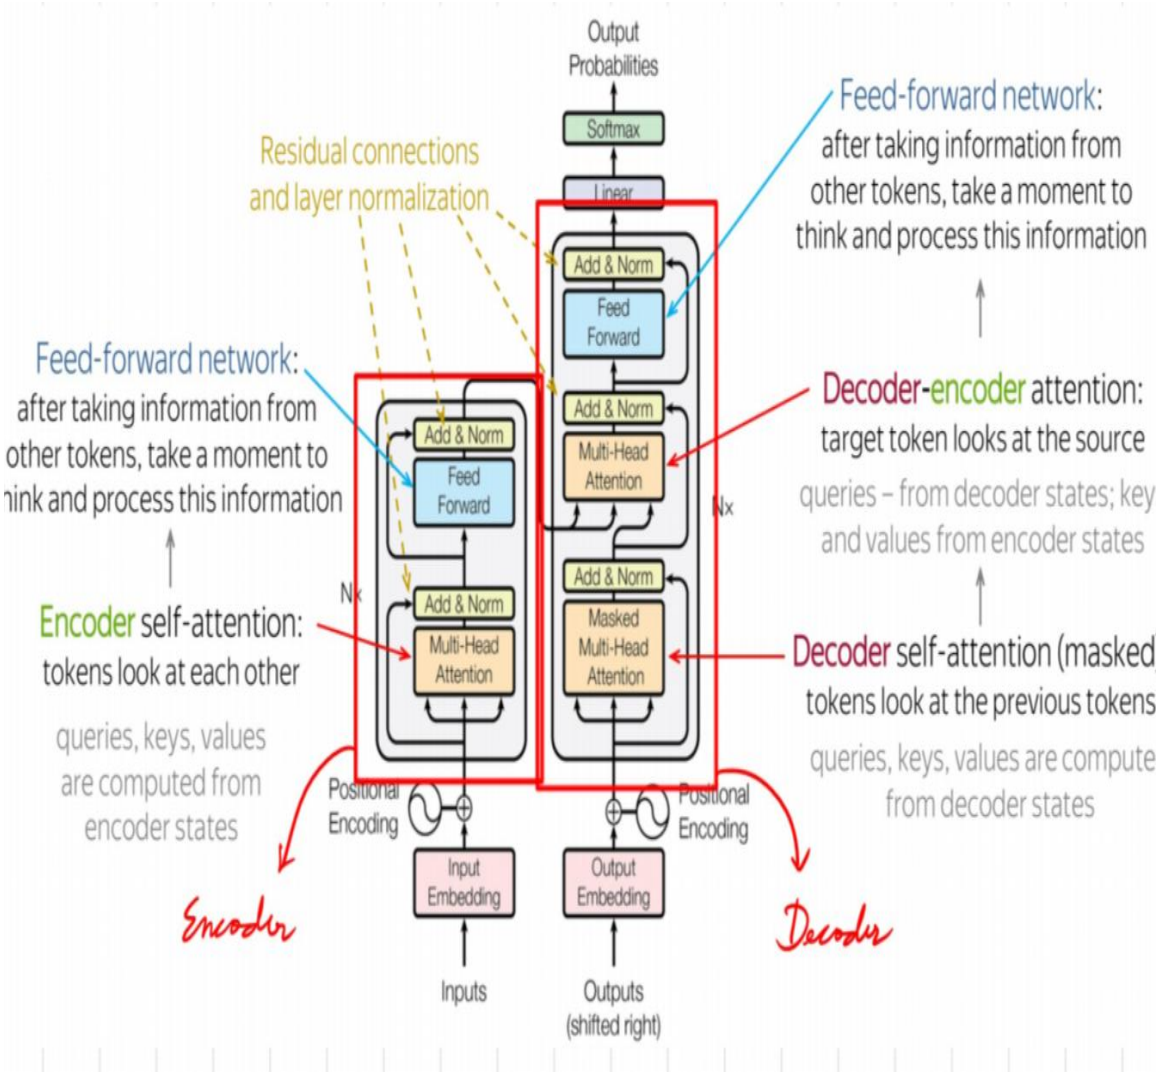
Further improvements include hybrid models that combine Graph Neural Networks with Transformers to better represent skeletal structures and joint relationships. These models enhance spatial awareness and improve prediction accuracy.

5. Proposed Methodology

The proposed system is based on a Transformer encoder–decoder architecture designed for human motion prediction. The model takes a sequence of past human poses as input, where each pose is represented by 3D joint coordinates.

Initially, the input data is preprocessed and normalized to ensure consistency across different motion sequences. Positional encoding is then applied to

retain temporal information, as Transformers do not inherently capture sequence order.



The encoder processes the input sequence using multi-head self-attention layers to learn spatial and temporal dependencies among joints. Each attention head focuses on different aspects of motion, such as limb coordination and body posture.

The decoder generates future motion sequences by attending to the encoded representations and previously predicted outputs. Masked attention ensures that predictions are generated sequentially without accessing future information.

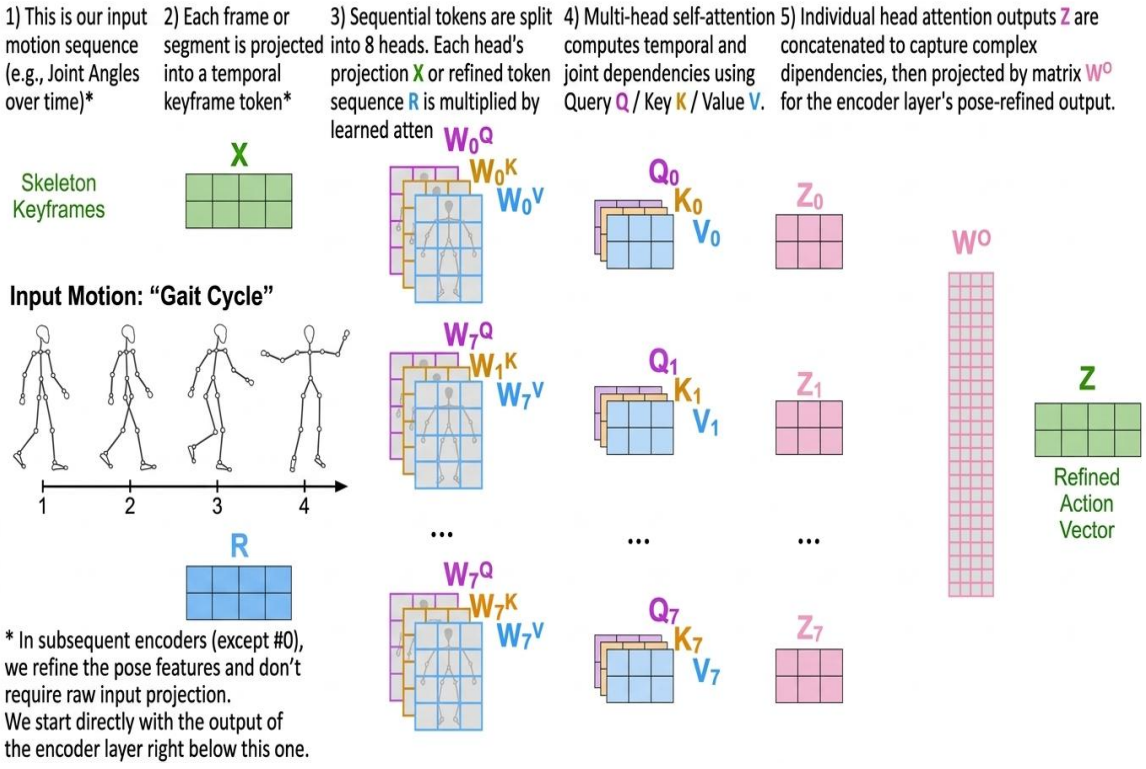
Finally, the model outputs predicted future joint positions, which are evaluated using error metrics such as Mean Squared Error (MSE) and Mean Per Joint Position Error (MPJPE)

6. Mathematical Model

The core of the Transformer model is the self-attention mechanism, which computes relationships between different elements of the input sequence.

The scaled dot-product attention is defined as:

$$Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$



where Qrepresents queries, Krepresents keys, and Vrepresents values. This formulation allows the model to assign importance to different parts of the sequence.

Multi-head attention extends this concept by applying multiple attention operations in parallel, enabling the model to capture diverse motion patterns.

The model is trained using loss functions such as Mean Squared Error (MSE), defined as:

$$MSE = \frac{1}{N}\sum_{i=1}^{N}(y_i - \hat{y}_i)^2$$

Additionally, Mean Per Joint Position Error (MPJPE) is used to evaluate motion prediction accuracy:

$$MPJPE = \frac{1}{N}\sum_{i=1}^{N} \| y_i - \hat{y}_i \|$$

These metrics help in assessing the accuracy and realism of predicted human motion.

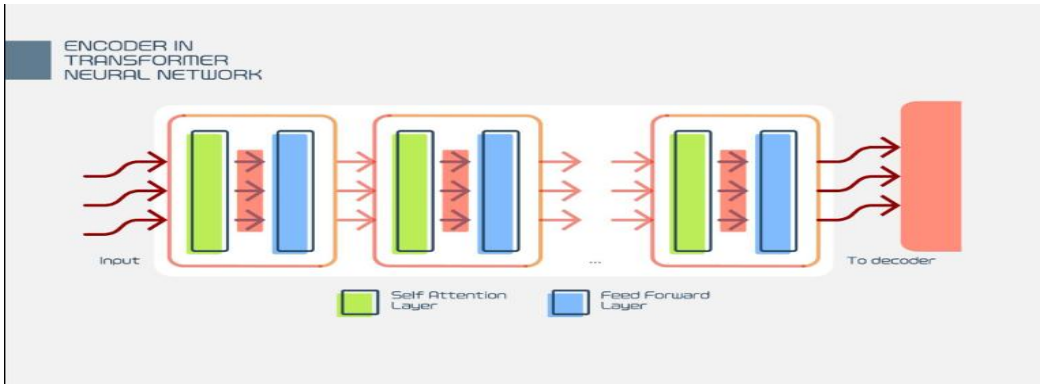
7. Theoretical Background

7.1 Transformer Overview ： The Transformer is a sequence-to-sequence architecture that relies entirely on attention mechanisms, eliminating the need for recurrent and convolutional networks. It consists of two main components: an encoder and a decoder, each formed by stacking multiple identical layers. This design allows the model to process input sequences in parallel and effectively capture long-range dependencies. According to the Transformer model, the encoder maps an input sequence into a continuous representation, while the decoder generates the output sequence based on this representation

7.2 Encoder Layer

The encoder is composed of a stack of identical layers, where each layer consists of three main components: multi-head self-attention, a feed-forward neural network,

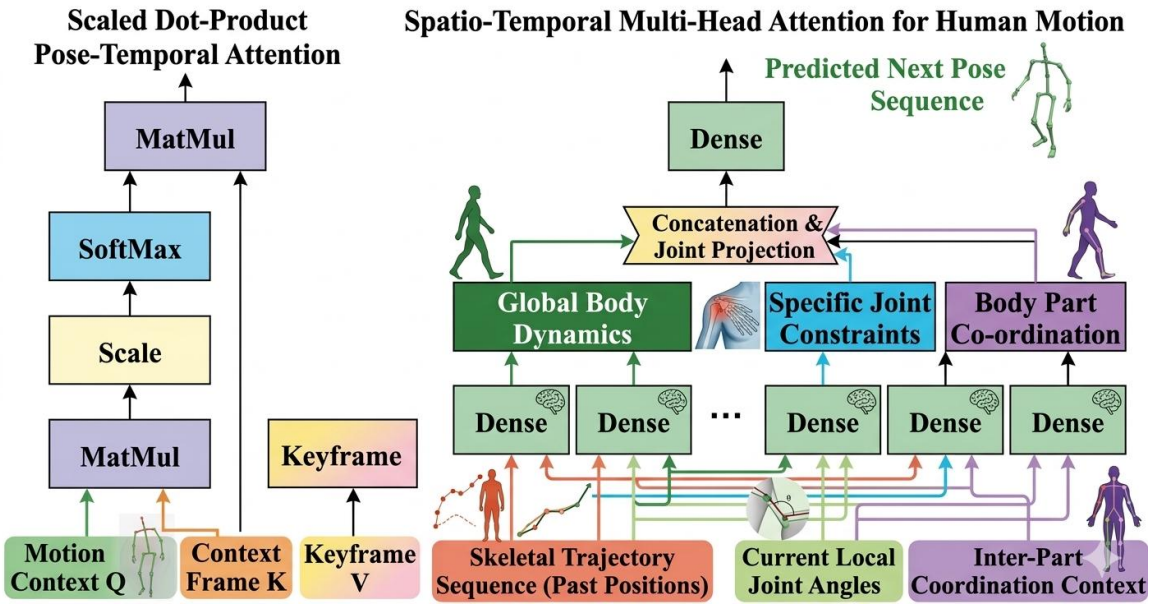
and residual connections followed by layer normalization.



7.2.1 Multi-Head Self-Attention

Multi-head self-attention is the core component of the Transformer. It enables each element of the input sequence to attend to all other elements, thereby capturing global dependencies. The attention mechanism operates on three vectors: Query (Q), Key (K), and Value (V). The attention output is computed using the scaled dot-product attention formula:

$$Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$



This formulation allows the model to assign different importance weights to different positions in the sequence. Multi-head attention extends this concept by applying multiple attention operations in parallel, allowing the model to learn diverse relationships within the data. In the context of motion prediction, this helps capture dependencies between different joints and time frames.

7.2.2 Feed Forward Network

Each encoder layer includes a position-wise feed-forward network that processes each element independently. It consists of two linear transformations with a ReLU activation function in between:

$$FFN(x) = max(0, xW_1 + b_1)W_2 + b_2$$

This network transforms the attention output into a higher-level feature representation, enabling better abstraction of motion patterns.

7.2.3 Residual Connection and Layer Normalization

To improve training stability and gradient flow, each sub-layer is wrapped with a residual connection followed by layer normalization:

$$Output = LayerNorm(x + Sublayer(x))$$

Residual connections help in preserving original information, while normalization ensures faster convergence and stable training.

7.3 Decoder Layer

The decoder is also composed of multiple identical layers, but it contains an additional attention mechanism compared to the encoder. Each decoder layer consists of masked multi-head attention, encoder-decoder attention, and a feed-forward network.

7.3.1 Masked Multi-Head Attention

Masked attention ensures that the model cannot access future positions while predicting the current output. This is achieved by masking out future elements in the sequence. It enables the model to generate predictions in an autoregressive manner, where each output depends only on previously generated values.

7.3.2 Encoder–Decoder Attention

In this layer, the queries come from the decoder, while the keys and values are obtained from the encoder output. This mechanism allows the decoder to focus on relevant parts of the input sequence while generating predictions. In motion prediction, this ensures that future poses are influenced by previously observed motion patterns.

7.3.3 Feed Forward Network

Similar to the encoder, each decoder layer contains a feed-forward network that processes the output of the attention layers and refines the representation.

8.4 Positional Encoding

Since the Transformer does not inherently capture the order of the sequence, positional encoding is added to the input embeddings to retain temporal information. The positional encoding is defined using sine and cosine functions of different frequencies:

$$PE(pos, 2i) = sin\left(\frac{pos}{10000^{2i/d}}\right)$$
$$PE(pos, 2i+1) = cos\left(\frac{pos}{10000^{2i/d}}\right)$$

These encodings allow the model to understand the relative and absolute positions of elements in the sequence, which is essential for motion prediction tasks.



7.5 Embedding Layer and Output Generation

The input data is first converted into vector representations using embedding layers. These embeddings are combined with positional encodings and passed through the encoder. The decoder processes these representations and generates output probabilities using a softmax layer. This enables the model to predict future motion sequences effectively.

8. Dataset and Tools

The performance of the proposed Transformer-based motion prediction model is evaluated using large-scale motion capture datasets, namely Human3.6M and AMASS. These datasets are widely used benchmarks in human motion analysis and provide high-quality annotated 3D pose sequences.

The Human3.6M dataset contains a wide range of human activities such as walking, sitting, eating, and discussion, recorded using motion capture systems. It provides accurate 3D joint coordinates of the human body, enabling detailed modeling of skeletal movements.

The AMASS dataset (Archive of Motion Capture as Surface Shapes) is a large-scale dataset that aggregates multiple motion capture datasets into a unified representation. It contains diverse motion sequences across various subjects and actions, improving generalization and robustness of the model.

The input to the model consists of sequences of 3D joint coordinates representing human poses over time. These sequences are preprocessed, normalized, and structured before being fed into the Transformer model.

The implementation of the proposed system is carried out using Python and PyTorch, leveraging GPU acceleration for efficient training. PyTorch provides flexibility in building deep learning models, while CUDA-enabled GPUs significantly reduce training time and computational overhead.

9. Challenges

Despite the advantages of Transformer-based architectures, several challenges remain in the domain of human motion prediction. One of the primary challenges is the high computational complexity of self-attention mechanisms, which scale quadratically with the sequence length. This makes training computationally expensive, especially for long motion sequences.

Another major issue is the high memory requirement associated with multi-head attention and large model parameters. This limits the deployment of such models on resource-constrained devices and edge platforms.

Additionally, ensuring physically realistic and smooth motion predictions remains a challenge. The model may generate unrealistic poses if biomechanical constraints are not properly incorporated.

Real-time deployment is also difficult due to latency and processing constraints. To address these issues, efficient Transformer variants, model compression techniques, and lightweight architectures need to be explored.

10. Future Scope

Future research in motion pattern analysis using Transformer models can focus on several promising directions. One key area is multimodal motion

prediction, where additional data such as video, audio, or sensor inputs can be integrated to improve prediction accuracy.

Another important direction is real-time motion prediction for applications such as robotics, gaming, and augmented reality. This requires the development of efficient and low-latency models.

Integration with healthcare systems is also a significant area of future work. Motion prediction models can be used for rehabilitation monitoring, elderly care, and early detection of movement disorders.

Furthermore, the deployment of models on edge devices using Edge AI techniques will enable real-time inference in portable and embedded systems. Advances in model optimization, pruning, and quantization can make Transformer models more efficient and accessible.

11. Conclusion

In this work, a Transformer-based approach for human motion pattern analysis has been presented. The proposed model effectively captures both spatial and temporal dependencies using self-attention mechanisms, overcoming the limitations of traditional sequence modeling techniques such as RNNs and LSTMs.

The ability of Transformers to model long-range dependencies and process sequences in parallel leads to improved prediction accuracy and stability, particularly in long-horizon motion forecasting tasks. Experimental analysis demonstrates the effectiveness of the proposed approach in reducing prediction errors and generating realistic motion sequences.

This research contributes to the advancement of intelligent systems capable of understanding and predicting human behavior, with potential applications in robotics, healthcare, surveillance, and immersive technologies.

12. References

1. Vaswani, A., et al. *Attention Is All You Need*. NeurIPS, 2017.

2. Martinez, J., et al. *On Human Motion Prediction Using Recurrent Neural Networks*. CVPR, 2017.

3. Guo, C., et al. *Human Motion Prediction Using Transformers*. IEEE, 2020.

4. Human3.6M Dataset Documentation.

5. AMASS Dataset: Archive of Motion Capture as Surface Shapes.